

MCIT 515

Fundamentals of Linear Algebra and Optimization Jean Gallier and Jocelyn Quaintance

Project 4: Hard Margin Support Vector Machine

The purpose of this project is to implement the hard margin support vector machine (version $h2$).

Recall that we would like to solve the following optimization problem:

Hard margin SVM (SVM_{h2}):

$$\begin{aligned} & \text{minimize} && \frac{1}{2} \|w\|^2 \\ & \text{subject to} && \\ & && w^\top u_i - b \geq 1 \quad i = 1, \dots, p \\ & && -w^\top v_j + b \geq 1 \quad j = 1, \dots, q, \end{aligned}$$

where $\{u_1, \dots, u_p\}$ is a set of p blue points and $\{v_1, \dots, v_q\}$ is a set of q red points in \mathbb{R}^n (here, $n = 2$). We assume that these two sets of points are separable. Figure 1 shows examples of points and separating planes, while figure ?? shows an example of the separating plane with the margins.




Figure 1: Examples of points and separating planes.

The problem is to find a separating hyperplane $H_{w,b}$ of equation $w^\top x - b = 0$ which maximizes the smallest distance δ from these data points, called the *margin*.

The margin is $\delta = 1/\|w\|$. The margin hyperplanes are the hyperplanes $H_{w,b+1}$ (the blue hyperplane) of equation $w^\top x - b - 1 = 0$ and $H_{w,b-1}$ (the red hyperplane) of equation




Figure 2: Examples of margins.

$w^\top x - b + 1 = 0$. In order to solve the above problem, we solve the dual program. The dual program is the following program:

Dual of the Hard margin SVM (SVM_{h2}):

$$\text{minimize} \quad \frac{1}{2} (\lambda^\top \quad \mu^\top) X^\top X \begin{pmatrix} \lambda \\ \mu \end{pmatrix} - (\lambda^\top \quad \mu^\top) \mathbf{1}_{p+q}$$

subject to

$$\sum_{i=1}^p \lambda_i - \sum_{j=1}^q \mu_j = 0$$

$$\lambda \geq 0, \mu \geq 0,$$

where X is the $n \times (p + q)$ matrix given by

$$X = \begin{pmatrix} -u_1 & \cdots & -u_p & v_1 & \cdots & v_q \end{pmatrix},$$

a matrix whose columns are the u_i and the v_j .

Then w is determined as follows:

$$w = -X \begin{pmatrix} \lambda \\ \mu \end{pmatrix} = \sum_{i=1}^p \lambda_i u_i - \sum_{j=1}^q \mu_j v_j.$$

To solve the dual using ADMM we need to determine the matrices P, q A and c . We have

$$P = X^\top X, \quad q = -\mathbf{1}_{p+q},$$

and since the only constraint is

$$\sum_{i=1}^p \lambda_i - \sum_{j=1}^q \mu_j = 0,$$

the matrix A is the $1 \times (p + q)$ row vector

$$A = (\mathbf{1}_p^\top \quad -\mathbf{1}_q^\top),$$

and the right-hand side c is the scalar

$$c = 0.$$

Obviously the matrix A has rank 1.

We obtain b using any i_0 such that $\lambda_{i_0} > 0$ and any j_0 such that $\mu_{j_0} > 0$. Since the corresponding constraints are active, we have

$$w^\top u_{i_0} - b = 1, \quad -w^\top v_{j_0} + b = 1,$$

so we obtain

$$b = w^\top (u_{i_0} + v_{j_0}) / 2.$$

For improved numerical stability, we can average over the sets of indices defined as $I_{\lambda>0} = \{i \in \{1, \dots, p\} \mid \lambda_i > 0\}$ and $I_{\mu>0} = \{j \in \{1, \dots, q\} \mid \mu_j > 0\}$. We obtain

$$b = w^\top \left(\left(\sum_{i \in I_{\lambda>0}} u_i \right) / |I_{\lambda>0}| + \left(\sum_{j \in I_{\mu>0}} v_j \right) / |I_{\mu>0}| \right) / 2.$$

(50 points) Write a Matlab program that implements the above method and solve for w and b . The ADMM optimizer is already implemented in `SVMhard2` in your project zip file.

The functions `buildhardSVM2`, `qsolve1`, `showdata`, `showSVMs2` and `makeline` (which is needed for `solveSVM2`) are given in the file folder `Matlabcode4`.

The parameter ρ is used by the function `qsolve1` that solve the optimization program using ADMM. We suggest using $\rho = 10$.

A difficulty that arises is to tune the tolerance parameters needed to deal with floating-point arithmetic. We suggested some tolerance parameters for you, but you are welcome to experiment.

Run your program on the following sets of data points:

```
v1 = [1 2 3 1 1 3 -1 -3;-1 0 -2 -0.5 -4 -3 -3 -3];
u1 = [-1 -1 0 1 -3 -4 0.5 3 0.5;0 1 2 3 0 -2 2 2.5 2.5];

u2 = 10.1*randn(2,20)+15;
v2 = -10.1*randn(2,20)-15;

u3 = 10.1*randn(2,20)+10;
v3 = -10.1*randn(2,20)-10;
```

```
u4 = 10.1*randn(2,50)+18;  
v4 = -10.1*randn(2,50)-18;
```

If you want to suppress printing `lamb` and `mu`, call the function using

```
[~, ~, w, b] = SVMhard2(rho,u,v)
```

What do you observe when you run hard SVM on the data set (`u3`, `v3`)? Write your answer in the report. Also try running these programs with several different random seeds and report the results. Try varying `rho` - what behavior do you notice? Are there choices of ρ where the program breaks?