

Project 9

In addition to answering the bolded questions on Coursera, also attach your notebook, both as `.ipynb` and `.html`.

Please follow the instructions below (e.g., setting random state values) to ensure that your answers match the solutions.

In this assignment, we will be using PennGrader, a Python package built by a former TA for autograding Python notebooks. PennGrader was developed to provide students with instant feedback on their answer. You can submit your answer and know whether it's right or wrong instantly. We then record your most recent answer in our backend database. You will have 100 attempts per test case, which should be more than sufficient.

NOTE: Please remember to remove the

```
raise NotImplementedError
```

after your implementation, otherwise the cell will not compile.

You currently work for a company that offers a subscription-based music streaming service. After an unsuccessful round of marketing efforts with low return on investments, your company is exploring other strategies to promote its services. Currently, its advertising strategy has been to bid on a set of keywords that someone on the marketing team put together, but the team has realized that they could improve on this strategy by refining the keywords based on different customer personas. Therefore, the company has asked you to help it identify a set of user segments so that the ads can be more customized.

For this task, you have been provided with the Music data which contains sociodemographic and music preference data on 4,914 users and your goal is to identify clusters of users based on this information.

Column	Description
Age	Age of the user
Gender	Gender of the user
Employment status	Employment status of the user
Annual income	Annual income of the user in USD
Usage per month	Average usage per month of the user measured in minutes
Top genre	The genre of music that is most streamed by the user
Num of days active	The number of days in the last 365 days that the user used the service

Getting Set Up

Meet our old friend - PennGrader! Fill in the cell below with your PennID and then run the following cell to initialize the grader.

Warning: Please make sure you only have one copy of the student notebook in your directory in Codio upon submission. The autograder looks for the variable `STUDENT_ID` across all notebooks, so if there is a duplicate notebook, it will fail.

```
In [1]: #PLEASE ENSURE YOUR STUDENT_ID IS ENTERED AS AN INT (NOT A STRING). IF NOT, TH  
E AUTOGRADER WON'T KNOW WHO  
#TO ASSIGN POINTS TO YOU IN OUR BACKEND
```

```
STUDENT_ID = 49731093          # YOUR 8-DIGIT PENNID GOES HERE  
STUDENT_NAME = "Newman Ilgenfritz" # YOUR FULL NAME GOES HERE
```

```
In [2]: import penngrader.grader  
  
grader = penngrader.grader.PennGrader(homework_id = 'ESE542_Online_Spring_2021  
_HW9', student_id = STUDENT_ID)
```

```
In [3]: # Let's import the relevant Python packages here  
# Feel free to import any other packages for this project  
  
#Data Wrangling  
import pandas as pd  
import numpy as np
```

Part A: Data Cleaning

Throughout this course, you have mostly been provided with relatively clean datasets but in the real world, data is usually a lot messier. This could be caused by data logging errors, bugs in the pipeline or even erroneous human input in surveys. Some common approaches of dealing with these observations are by dropping them, replacing the erroneous values with the mean/median, winsorizing, etc. Here, we will just drop them.

1. Load the Music data and drop any rows with `NaN/null` values. Assign the dataframe to the variable `music`.

```
In [4]: music = pd.read_csv('music.csv')

print('music.shape before dropna = ', music.shape, '\n')

# Drop all rows with NaN/null inplace

music.dropna(axis=0, how='any', thresh=None, subset=None, inplace=False)

# df.reset_index(drop=True)
music.reset_index(drop=True)
print('music.head: ')
print(music.head(), '\n')
print('music.tail: ')
print(music.tail(), '\n')
print('music.shape = ', music.shape, '\n')

#raise NotImplementedError
```

music.shape before dropna = (4914, 7)

music.head:

	age	gender	employment_status	annual_income	usage_per_month	top_genre
0	20	Male	Employed full-time	115000	847	Rock
1	22	Male	Employed full-time	42100	256	Rock
2	18	Male	Employed full-time	114750	1232	Country
3	19	Female	Employed full-time	118250	2310	Rock
4	33	Male	Employed full-time	111450	568	Pop

	num_of_days_active
0	210
1	222
2	161
3	176
4	185

music.tail:

	age	gender	employment_status	annual_income	usage_per_month	\
4909	31	Female	Unemployed	17150	1291	
4910	41	Male	Employed full-time	19150	905	
4911	37	Male	Employed full-time	39000	535	
4912	43	Female	Student	138800	13	
4913	44	Male	Employed full-time	27000	1184	

	top_genre	num_of_days_active
4909	Rock	220
4910	Pop	166
4911	Rock	203
4912	Rock	126
4913	Pop	230

music.shape = (4914, 7)

```
In [5]: grader.grade(test_case_id = 'test_read_music', answer = music.shape)
```

Correct! You earned 1/1 points. You are a star!

Your submission has been successfully recorded in the gradebook.

1. Plot bar charts and histograms of the data to visualize the distributions. Does anything seem unusual? Good to think about: can you encapsulate plotting data into a function?
 - Are there more male or female users? Store your answer in `male_or_female` ('M'/'F').
 - Which genre is the most popular among users? Store your answer in `pop_genre`.
 - How many users are students? Store your answer in `student_num`.

```
In [6]: from matplotlib import pyplot as plt
%matplotlib inline

#music.plot.bar()
#plt.title("Magnitude of each data column:")
#plt.xlabel("Column type")
#plt.ylabel("Magnitude of each")
#plt.show()
#plt.close()

#music.plot.hist()
#plt.show()
#plt.close()

mpt = music.pivot_table(index=['gender'], aggfunc='size')
print('mpt: ')
print(mpt, '\n')

mpt = music.pivot_table(index=['top_genre'], aggfunc='size')
print('mpt: ')
print(mpt, '\n')

mpt = music.pivot_table(index=['employment_status'], aggfunc='size')
print('mpt: ')
print(mpt, '\n')
```

```
mpt:
gender
Female          2350
Male            2482
Other              22
Prefer not to say    60
dtype: int64
```

```
mpt:
top_genre
Classical      721
Country        556
Hip Hop         99
Indie          289
Jazz           265
Pop           2149
Rock           835
dtype: int64
```

```
mpt:
employment_status
Employed full-time    3037
Employed part-time    513
Student              1000
Unemployed            364
dtype: int64
```

```
In [7]: male_or_female = 'M'
```

```
In [8]: grader.grade(test_case_id = 'test_visual1', answer = male_or_female)
```

Correct! You earned 1/1 points. You are a star!

Your submission has been successfully recorded in the gradebook.

```
In [9]: pop_genre = 'pop'
```

```
In [10]: grader.grade(test_case_id = 'test_visual2', answer = pop_genre)
```

Correct! You earned 1/1 points. You are a star!

Your submission has been successfully recorded in the gradebook.

```
In [11]: student_num = 1000
```

```
In [12]: grader.grade(test_case_id = 'test_visual3', answer = student_num)
```

Correct! You earned 1/1 points. You are a star!

Your submission has been successfully recorded in the gradebook.

1. Take a closer look at 'age' and 'num of days active'. Does anything look peculiar? Store the number of invalid 'age' in `bad_age_num`, invalid number of 'num of days active' in `bad_active_num`. Drop these rows and also assign the new dataframe to the variable `music`. *Hint*: Recall the values that age and number of days active can take on.

```
In [13]: # Enter your code here
mpt = music.pivot_table(index=['age'], aggfunc='size')
print('mpt: ')
print(mpt, '\n')

baNum = [-29, -27, -21, -10, -1]
bad_age_num = len(baNum) # 5
print('bad_age_num: ', bad_age_num, '\n')
```

mpt:

age

-29	1
-27	1
-21	1
-10	1
-1	1
16	133
17	128
18	144
19	146
20	139
21	144
22	147
23	152
24	142
25	140
26	129
27	131
28	142
29	124
30	165
31	154
32	164
33	187
34	160
35	158
36	163
37	130
38	167
39	160
40	145
41	118
42	118
43	117
44	119
45	112
46	112
47	124
48	107
49	106
50	107
51	4
52	4
53	6
54	5
55	3
56	4
57	8
58	5
59	5
60	8
61	6
62	2
63	6
64	7
65	2


```
dtype: int64
```

```
bad_age_num: 5
```

```
In [14]: grader.grade(test_case_id = 'test_drop1', answer = bad_age_num)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

```
In [15]: # Enter your code here
mpt = music.pivot_table(index=['num_of_days_active'], aggfunc='size')
print('mpt: ')
#print(mpt.head(50), '\n')
print(mpt.tail(5), '\n')

baActNum = [380, 392]
bad_active_num = len(baActNum) # 2
print('bad_active_num: ', bad_active_num, '\n')
```

```
mpt:
num_of_days_active
329      1
346      1
350      1
380      1
392      1
dtype: int64
```

```
bad_active_num: 2
```

```
In [16]: grader.grade(test_case_id = 'test_drop2', answer = bad_active_num)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

```
In [17]: # Drop rows with invalid entries in-place

music = music[music.age > 1]
mpt = music.pivot_table(index=['age'], aggfunc='size')
print('mpt: ')
print(mpt, '\n')

music = music[music.num_of_days_active <= 365]
mpt = music.pivot_table(index=['num_of_days_active'], aggfunc='size')
print('mpt: ')
print(mpt, '\n')

print('\nmusic.shape = ', music.shape, '\n')
```

```
mpt:
age
16    133
17    128
18    144
19    146
20    139
21    144
22    147
23    152
24    142
25    140
26    129
27    131
28    142
29    124
30    165
31    154
32    164
33    187
34    160
35    158
36    163
37    130
38    167
39    160
40    145
41    118
42    118
43    117
44    119
45    112
46    112
47    124
48    107
49    106
50    107
51     4
52     4
53     6
54     5
55     3
56     4
57     8
58     5
59     5
60     8
61     6
62     2
63     6
64     7
65     2
dtype: int64
```

```
mpt:
num_of_days_active
31     1
```

```
34      1
41      1
42      1
45      1
      ..
323     1
326     1
329     1
346     1
350     1
Length: 262, dtype: int64
```

```
music.shape = (4907, 7)
```

```
In [18]: grader.grade(test_case_id = 'test_drop3', answer = music.shape)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

Part B: Clustering

Use K-means clustering to find a set of user groups using the following features as inputs: 'age', 'annual income', 'usage per month' and 'number of days active'.

1. Standardize the data. *Hint:* `sklearn.preprocessing.scale` may be helpful.

```
In [19]: from sklearn import preprocessing
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import scale

musicQuant = music.drop(['gender', 'employment_status', 'top_genre'], axis=1)
print('musicQuant.head: ')
print(musicQuant.head(), '\n')

musicQuantScld = preprocessing.scale(musicQuant, axis=0, with_mean=True, with_
std=True, copy=True)
print('musicQuantScld: ')
print(musicQuantScld, '\n')

X_scaled = musicQuantScld
```

```
musicQuant.head:
   age  annual_income  usage_per_month  num_of_days_active
0   20         115000             847             210
1   22          42100             256             222
2   18         114750            1232             161
3   19         118250            2310             176
4   33         111450             568             185
```

```
musicQuantScld:
[[-1.26120736  1.70825627 -0.32086458  0.62887233]
 [-1.06361331  0.02875876 -0.94875035  0.89592062]
 [-1.45880141  1.70249667  0.08816421 -0.46157486]
 ...
 [ 0.41834209 -0.0426602  -0.65233727  0.47309416]
 [ 1.01112424  2.25656958 -1.20691657 -1.24046571]
 [ 1.10992127 -0.31912069  0.03716841  1.07395282]]
```

```
In [20]: grader.grade(test_case_id = 'test_scale', answer = X_scaled)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

1. Calculate the sum of squared distances of observations to their closest cluster center for $K \in [1, 10]$ and add these values to a list called `ssd`. Use the default hyperparameters and set `random_state=42`. Store the minimum value of `ssd` in `min_ssd`. Comment on your thoughts: shall we choose the model with minimum `ssd`?

Hint: You may want to write a for-loop and refer to the `sklearn.cluster.KMeans` (<https://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>) documentation (read through the Attributes section).

```
In [21]: clusters = 3
maxIter = 300
kmeansArgs = {"init": "random", "n_init": 10, "max_iter": maxIter, "random_state": 42}

ssd = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, **kmeansArgs)
    kmeans.fit(X_scaled)
    ssd.append(kmeans.inertia_)

print('ssd: ')
print(ssd, '\n')

min_ssd = min(ssd)
print('min_ssd: ', min_ssd, '\n')
```

```
ssd:
[19628.000000000004, 15279.939975664749, 12077.981652729, 10202.025406603878,
8889.815844988263, 8021.102051116282, 7353.001733920524, 6775.933069916614, 6
325.373776598275, 6014.832512563668]
```

```
min_ssd: 6014.832512563668
```

```
In [22]: grader.grade(test_case_id = 'test_min_ssd', answer = min_ssd)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

1. Plot the values in `ssd` against the number of clusters, K for $K \in [1, 10]$.

```
In [23]: x = np.arange(1, 11)
y = ssd
plt.style.use("fivethirtyeight")
plt.plot(x, y)
plt.xticks(range(1, 11))
plt.xlabel("Number of Clusters")
plt.ylabel("SSE")
plt.show()
plt.close()
```



- Based on this plot, what is the best number of clusters to set using the Elbow Method? The Elbow Method is the point where diminishing returns are no longer worth the additional cost. It will be the point before which the later points taper off or increase by very little. (More information on the Elbow Method can be found [here](https://en.wikipedia.org/wiki/Elbow_method_(clustering)) ([https://en.wikipedia.org/wiki/Elbow_method_\(clustering\)](https://en.wikipedia.org/wiki/Elbow_method_(clustering)))#~:text=In%20cluster%20analysis%2C%20the%20Assign this value to the variable `optimal_K` (There might be multiple optimal k values, just enter one of those).

```
In [25]: #import kneed
#from kneed import KneeLocator

#kl = KneeLocator( range(1, 11), ssd, curve="convex", direction="decreasing" )

# optimal_K = kl.elbow
optimal_K = 4
```

```
In [26]: grader.grade(test_case_id = 'test_best_K', answer = optimal_K)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

1. Retrain the K-means clustering algorithm with `n_clusters=optimal_K` . This will be your final model. Name your model `best_kmeans` and set `random_state=42` .

```
In [27]: clusters = optimal_K
maxIter = 300
kmeansArgs = {"init": "random", "n_init": 10, "max_iter": maxIter, "random_state": 42}
kmeans = KMeans(n_clusters=clusters, **kmeansArgs)
#kmeans.fit(X_scaled)

best_kmeans = kmeans.fit(X_scaled)

print('best_kmeans.cluster_centers_ :')
print(best_kmeans.cluster_centers_, '\n')

best_kmeans.cluster_centers_ :
[[ 9.54499771e-01 -3.96147495e-01 -4.38869616e-01 -1.05320102e-01]
 [-8.88934157e-01 -4.13304281e-01 -7.75288958e-01 -1.44681625e-03]
 [ 4.47434260e-01  1.94395051e+00  2.30392334e-01 -1.58390715e-02]
 [-7.08953631e-02 -4.08179560e-01  1.17531767e+00  1.05911249e-01]]
```

```
In [28]: grader.grade(test_case_id = 'test_best_model', answer = (optimal_K, best_kmeans.cluster_centers_))
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

Part C: Visualizing

As a data scientist, it is important to be able to translate your findings to colleagues who may not have the same level of technical knowledge as you do; visualizations help a lot! Therefore, you have decided to plot the clusters. However, since there are four dimensions that you want to plot – age, annual income, usage per month and number of days active – you will need to utilize dimensionality reduction techniques to be able to plot this on a 2-D plane.

1. Train another K-means model using the following features as inputs: age , annual income , usage per month and number of days active . Fit your model on `X_scaled` , name your predicted result as `pred` . Use `n_clusters=2` and `random_state=42` .

```
In [ ]: # Enter your code here

raise NotImplementedError
```

```
In [ ]: grader.grade(test_case_id = 'test_kmeans_plot', answer = pred)
```


1. Find the first **two** principal components of the scaled data and save the new data to variable `X_reduced` . Name your model `pca` .

Hint: Use `sklearn.decomposition.PCA` with `random_state=42` .

```
In [ ]: pca =  
  
raise NotImplementedError
```

```
In [ ]: grader.grade(test_case_id = 'test_pca', answer = (pca.n_features_, pca.n_components_))
```

1. Plot a scatterplot with the 1st principal component on the x -axis and the 2nd principal component on the y -axis. Color each point by the cluster that it is in from (1). Comment on your observations.

Hint: Check out `seaborn.FacetGrid` . To use it, you will need your principal components and predicted clusters in one dataframe.

```
In [ ]: # Enter your code here  
  
raise NotImplementedError
```