ESE 542: Statistics for Data Science

Chapter 3 Project

In addition to answering the **bolded** questions on Coursera, also attach your notebook, both as `.ipynb` and `.html`.

In the following exercise, we will perform linear regression to fit various data sets and to predict outputs. Perform the following analyses by starting a new notebook.

1 Part A

First, we will use the Ch3PartA dataset to generate polynomial regressions using `scikitlearn`. This dataset contains 100 observations of points x and their corresponding response, y. The data is divided into a training set (x_{tr}, y_{tr}) and a test set (x_{te}, y_{te}).

1. Load `Ch3PartA.csv` into your notebook.
2. Create a scatter plot of: (a) y_{tr} against x_{tr} and another of (b) y_{te} against x_{te}. Notice the similarities and differences between the plots. **Answer on Coursera: What is the maximum value of y in the training set? In the test set?**
3. Generate the necessary features to fit polynomial regressions up to the 20^{th} degree (up to and including the x^{20} term) on the training data. *Hint*: You will be fitting multivariate linear regression models with polynomial features of x. Familiarize yourself with `sklearn.preprocessing.PolynomialFeatures`.
4. Calculate the training MSE and the test MSE for 20 polynomial models up to degree 20. Store these as lists named `mse_train` and `mse_test`respectively. *Hint*: Familiarize yourself with `sklearn.metrics.mean_squared_error` and try to automate the process.
5. Generate a plot of both the training MSE and test MSE against flexibility for degrees 1 to 20. **Answer on Coursera: What is the minimum training MSE? Test MSE?**
6. From your plot, make an educated guess about the polynomial degree of the function that was used to generate the data. Then, give an estimate of the irreducible error $\text{Var}(\epsilon)$ for the

optimal model on both the training set and test set. *Hint*: Revisit the section on hypothesis testing. Think about the relationship between MSE, RSS, and RSE. **Enter your answers on Coursera.**

2 Part B

Next, we will use the Ch3PartB dataset to observe the effects of collinearity using `statsmodels`. This dataset contains 100 observations of points x_1, x_2, and y, the response variable.

1. Load the data from `Ch3PartB.csv` into a `pandas DataFrame`.
2. Show a scatterplot displaying the relationship between x_1 and x_2? **Answer on Coursera: What is the correlation coefficient between x_1 and x_2?**
3. Using the data, fit a least squares regression to predict y using x_1 and x_2. Describe your results. *Hint*: Familiarize yourself with `statsmodels.formula.api.ols`. **Answer on Coursera: What are the estimates $\hat{\beta}_0$, $\hat{\beta}_1$ and $\hat{\beta}_2$? At a 95% confidence level, can you reject the null hypothesis $H_0 : \beta_1 = 0$? What about $H_0 : \beta_2 = 0$?**
4. Now fit a least squares regression to predict y using only x_1. Comment on your results. **Answer on Coursera: Can you reject the null hypothesis $H_0 : \beta_1 = 0$?**
5. Fit a least squares regression to predict y using only x_2. Comment on your results. **Answer on Coursera: Can you reject the null hypothesis $H_0 : \beta_2 = 0$?**
6. Do Part B Questions 3-5 contradict each other? Explain why.