

Project 4

In addition to answering the bolded questions on Coursera, also attach your notebook, both as `.ipynb` and `.html`.

This project should be answered using the `Weekly` data set (attached). This data contains 1,089 weekly stock market percentage returns for 21 years, from the beginning of 1990 to the end of 2010.

Details about the columns in the data are summarized below:

- `Year` : The year that the observation was recorded
- `Lag1` : Percentage return for previous week
- `Lag2` : Percentage return for 2 weeks previous
- `Lag3` : Percentage return for 3 weeks previous
- `Lag4` : Percentage return for 4 weeks previous
- `Lag5` : Percentage return for 5 weeks previous
- `Volume` : Volume of shares traded (average number of daily shares traded in billions)
- `Today` : Percentage return for this week
- `Direction` : A factor with levels Down and Up indicating whether the market had a positive or negative return on a given week

In this assignment, we will be using PennGrader, a Python package built by a former TA for autograding Python notebooks. PennGrader was developed to provide students with instant feedback on their answer. You can submit your answer and know whether it's right or wrong instantly. We then record your most recent answer in our backend database. You will have 100 attempts per test case, which should be more than sufficient.

NOTE : Please remember to remove the

```
raise NotImplementedError
```

after your implementation, otherwise the cell will not compile.

Getting Setup

Please run the below cells to get setup with the autograder. If you need to install packages, please copy these lines into the Terminal!

```
In [70]: # !pip install pandas==1.0.5 --user
# pip install penngrader --user
```

```
In [72]: # pip install seaborn --user
# pip install scikit-learn --user
# pip install statsmodels --user
```

Let's try PennGrader out! Fill in the cell below with your PennID and then run the following cell to initialize the grader.

Warning: Please make sure you only have one copy of the student notebook in your directory in Codio upon submission. The autograder looks for the variable `STUDENT_ID` across all notebooks, so if there is a duplicate notebook, it will fail.

```
In [73]: #PLEASE ENSURE YOUR STUDENT_ID IS ENTERED AS AN INT (NOT A STRING). IF
#NOT, THE AUTOGRADER WON'T KNOW WHO
#TO ASSIGN POINTS TO YOU IN OUR BACKEND

STUDENT_ID = 49731093 # YOUR 8-DIGIT PENNID GOES HERE
STUDENT_NAME = "Newman Ilgenfritz" # YOUR FULL NAME GOES HERE
```

```
In [74]: import penngrader.grader

grader = penngrader.grader.PennGrader(homework_id = 'ESE542_Online_Spring_2021_HW4', student_id = STUDENT_ID)
```

```
In [75]: # Let's import the relevant Python packages here
```

```
# Feel free to import any other packages for this project

# Data Wrangling
import pandas as pd
import numpy as np

# Statistics
import statsmodels.formula.api as smf
import statsmodels.api as sm

# Plotting
import matplotlib.pyplot as plt

%matplotlib inline
```

We're also going to run a quick (0-point) check that the pandas version set up here is correct. If you fail this, please open a Terminal window and run `pip install pandas==1.0.5 --user` otherwise none of the tests will work!

```
In [76]: grader.grade(test_case_id = 'A0_pandas_test', answer = str(pd.__version__))
```

Correct! You earned 0/0 points. You are a star!

Your submission has been successfully recorded in the gradebook.

Part A

We are first interested in trying to predict the direction of the returns.

To start, load `Weekly.csv` into your notebook.

```
In [77]: weeklyFile = pd.read_csv('Weekly.csv').copy()
weekly= pd.DataFrame(weeklyFile)
```

```
In [78]: grader.grade(test_case_id = 'A0_weekly_test', answer = weekly)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

A1.

First, transform our `Direction` variable into a numerical feature that is equal to 1 if `Direction = Up`. Then, pass the dataframe into the test case to make sure it's working properly!

```
In [79]: #raise NotImplementedError
encode = lambda x: 1 if x == 'Up' else 0
weekly['Direction'] = weekly['Direction'].map(encode)
print('weekly after encoding "direction": ')
print(weekly.head(), '\n')
```

weekly after encoding "direction":

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
0	1990	0.816	1.572	-3.936	-0.229	-3.484	0.154976	-0.270	0
1	1990	-0.270	0.816	1.572	-3.936	-0.229	0.148574	-2.576	0
2	1990	-2.576	-0.270	0.816	1.572	-3.936	0.159837	3.514	1
3	1990	3.514	-2.576	-0.270	0.816	1.572	0.161630	0.712	1
4	1990	0.712	3.514	-2.576	-0.270	0.816	0.153728	1.178	1

```
In [80]: grader.grade(test_case_id = 'A1_direction_test', answer = weekly)
```

Correct! You earned 1/1 points. You are a star!

Your submission has been successfully recorded in the gradebook.

Produce some numerical and graphical summaries of the `Weekly` data. Do there appear to be any patterns?

```
In [81]: #raise NotImplementedError
```

```

print('features correlations: ')
print(weekly.corr(), '\n')
print('\nMatrix of pairwise scatter plots:')
axes = pd.plotting.scatter_matrix(weekly, alpha=0.30, color="brown")
#pd.plotting.scatter_matrix(weekly, alpha=0.30, color="brown")
#g = pd.plotting.scatter_matrix(weekly)
#plt.show()
#plt.close()
#sns.pairplot(weekly)
plt.show()
plt.close()
print()

```


features correlations:

	Year	Lag1	Lag2	Lag3	Lag4	Lag5
\						
Year	1.000000	-0.032289	-0.033390	-0.030006	-0.031128	-0.030519
Lag1	-0.032289	1.000000	-0.074853	0.058636	-0.071274	-0.008183
Lag2	-0.033390	-0.074853	1.000000	-0.075721	0.058382	-0.072499
Lag3	-0.030006	0.058636	-0.075721	1.000000	-0.075396	0.060657
Lag4	-0.031128	-0.071274	0.058382	-0.075396	1.000000	-0.075675
Lag5	-0.030519	-0.008183	-0.072499	0.060657	-0.075675	1.000000
Volume	0.841942	-0.064951	-0.085513	-0.069288	-0.061075	-0.058517
Today	-0.032460	-0.075032	0.059167	-0.071244	-0.007826	0.011013
Direction	-0.022200	-0.050004	0.072696	-0.022913	-0.020549	-0.018168

	Volume	Today	Direction
Year	0.841942	-0.032460	-0.022200
Lag1	-0.064951	-0.075032	-0.050004
Lag2	-0.085513	0.059167	0.072696
Lag3	-0.069288	-0.071244	-0.022913
Lag4	-0.061075	-0.007826	-0.020549
Lag5	-0.058517	0.011013	-0.018168
Volume	1.000000	-0.033078	-0.017995
Today	-0.033078	1.000000	0.720025
Direction	-0.017995	0.720025	1.000000

Matrix of pairwise scatter plots:





Include a brief description of what relationships and correlations you find.

```
In [82]: relationships = '''
        the correlations between the lag variables and today's returns are
        close to zero.
        The only significant positive correlations are between Year and Vol
        ume (0.841942),
        and between Direction and Today (0.720025).

        ...

        #raise NotImplementedError
```

```
In [83]: grader.grade(test_case_id = 'A1_relationships_test', answer = relations
        hips)
```

Correct! You earned 1/1 points. You are a star!

Your submission has been successfully recorded in the gradebook.

A2.

Use the full data set to perform a logistic regression with `Direction` as the response and the five lag variables as predictors.

```
In [84]: #raise NotImplementedError
fit1 = smf.glm('Direction~Lag1+Lag2+Lag3+Lag4+Lag5', data = weekly, fam
ily=sm.families.Binomial()).fit()
print(fit1.summary(), '\n')
print('Predicted probability of going Up:', fit1.predict(), '\n') #Predict
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Direction   No. Observations:
1089
Model:                  GLM        Df Residuals:
1083
Model Family:          Binomial    Df Model:
5
Link Function:         logit       Scale:
1.0000
Method:                IRLS       Log-Likelihood:
-743.37
Date:                  Mon, 15 Feb 2021   Deviance:
1486.7
Time:                  17:17:10   Pearson chi2:
1.09e+03
No. Iterations:        4
Covariance Type:      nonrobust
=====
=====
               coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
```

Intercept	0.2303	0.062	3.712	0.000	0.109
0.352					
Lag1	-0.0401	0.026	-1.522	0.128	-0.092
0.012					
Lag2	0.0602	0.027	2.249	0.025	0.008
0.113					
Lag3	-0.0151	0.027	-0.566	0.571	-0.067
0.037					
Lag4	-0.0268	0.026	-1.013	0.311	-0.079
0.025					
Lag5	-0.0135	0.026	-0.512	0.609	-0.065
0.038					

=====

=====

Predicted probability of going Up: [0.59979417 0.5926544 0.57835811
... 0.59836445 0.56799844 0.52963133]

◀ ▶

Pass in the regression equation to `logit_equation` below. Hint: You do not need the coefficients of the equation yet, just which variables you want to include in the model. Your answer should look something like `Response~Var1+Var2` which is the input for `statsmodels.formula.api`.

In [85]: `logit_equation = 'Direction~Lag1+Lag2+Lag3+Lag4+Lag5'`

In [86]: `grader.grade(test_case_id = 'A2_logit_test', answer = logit_equation)`

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

A3.

Use the `summary()` function to print the results. Do any of the predictors appear to be statistically significant? Which predictors appear to be statistically significant?

```
In [87]: #raise NotImplementedError
print(fit1.summary(), '\n')
print('Predicted probability of going Up:', fit1.predict(), '\n') #Predict
```

```
Generalized Linear Model Regression Results
=====
=====
Dep. Variable:                Direction    No. Observations:
1089
Model:                        GLM          Df Residuals:
1083
Model Family:                Binomial      Df Model:
5
Link Function:                logit         Scale:
1.0000
Method:                       IRLS         Log-Likelihood:
-743.37
Date:                         Mon, 15 Feb 2021    Deviance:
1486.7
Time:                         17:17:15          Pearson chi2:
1.09e+03
No. Iterations:                4
Covariance Type:              nonrobust
=====
=====
                                coef    std err          z      P>|z|      [0.025
0.975]
-----
-----
Intercept          0.2303      0.062      3.712      0.000      0.109
0.352
Lag1              -0.0401      0.026     -1.522      0.128     -0.092
0.012
Lag2               0.0602      0.027      2.249      0.025      0.008
0.113
```

Lag3	-0.0151	0.027	-0.566	0.571	-0.067
0.037					
Lag4	-0.0268	0.026	-1.013	0.311	-0.079
0.025					
Lag5	-0.0135	0.026	-0.512	0.609	-0.065
0.038					

=====

=====

Predicted probability of going Up: [0.59979417 0.5926544 0.57835811
... 0.59836445 0.56799844 0.52963133]

◀ ▶

Type the number of apparently significant variables into `num_significant` and the names of the variables into the list `var_significant` -- the test case will only give points if both variables are correct!

```
In [88]: # 0.05 > Lag2 therefore reject null for Lag2 only
# print(fit1.pvalues) # same as p values in summary, but carried out to
# more sig digits

num_significant = 1
var_significant = ['Lag2'] # This should be a list!
# raise NotImplementedError
```

```
In [89]: grader.grade(test_case_id = 'A3_significant_test', answer = (num_signif
icant, var_significant))
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

A4.

Compute the overall fraction of correct predictions. Name this variable `fraction_correct_all`. What is the overall fraction of correct predictions?

```
In [90]: print('\nnumber of obs: ', len(weekly))
encode = lambda x: 1 if x > 0.5 else 0
predictions = pd.DataFrame({'Direction': [encode(x) for x in fit1.predict()]})
print('predictions.head(): ')
print(predictions.head(), '\n')
from sklearn.metrics import confusion_matrix
x = confusion_matrix(weekly['Direction'], predictions['Direction'])
print('confusion_matrix: ')
print(x, '\n')
from sklearn.metrics import accuracy_score
# print('The accuracy score is', 100 * accuracy_score(data['Direction'], predictions['Direction']), '%')
aScore = accuracy_score(weekly['Direction'], predictions['Direction'])
print('aScore: ', aScore)
fraction_correct_all = aScore
```

number of obs: 1089

predictions.head():

	Direction
0	1
1	1
2	1
3	0
4	1

confusion_matrix:

```
[[ 49 435]
 [ 41 564]]
```

aScore: 0.5629017447199265

```
In [91]: print(f'Overall fraction of correct predictions is {fraction_correct_all}')
Overall fraction of correct predictions is 0.5629017447199265
```

```
In [92]: grader.grade(test_case_id = 'A4_fraction_test', answer = fraction_correct_all)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

A5.

Now fit the logistic regression model using a training data period from 1990 to 2007, with `Lag2` as the only predictor.

Compute the overall fraction of correct predictions for the held out data (that is, the data from 2008, 2009 and 2010) and assign it to a variable called `fraction_correct_test`. What is the overall fraction of correct predictions?

```
In [93]: from sklearn.metrics import confusion_matrix
#Train and test split
train = weekly[weekly["Year"] < 2008]
print('train.head: ')
print(train.head(), '\n')
test = weekly[weekly["Year"] >= 2008]
print('test.head: ')
print(test.head(), '\n')
print('test obs: ', len(test), '\n')

# Model
logit_equation = 'Direction~Lag2'
fit2 = smf.glm(logit_equation, data = train, family=sm.families.Binomial()).fit()
print(fit2.summary(), '\n')
#print('Predicted probability of going Up:', fit1.predict(), '\n') #Predict

Xtest = test[weekly.columns[1:-2]]
print('Xtest.head: ')
```

```

print(Xtest.head(), '\n')
ytest = test[weekly.columns[-1]]
print('ytest.head: ')
print(ytest.head(), '\n')
#yprobs = fit2.predict(Xtest)
yprobs = fit2.predict(Xtest)

ypred = [0 if y < 0.5 else 1 for y in yprobs]
x = confusion_matrix(ytest, ypred)
print('cm: ')
print(x, '\n')

from sklearn.metrics import accuracy_score
aScore = accuracy_score(ytest, ypred)

print('aScore: ', aScore)
#print('The accuracy score is', 100 * accuracy_score(data['Direction'],
predictions['Direction']), '%')
fraction_correct_test = aScore

```

```

#raise NotImplementedError

```

train.head:

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
0	1990	0.816	1.572	-3.936	-0.229	-3.484	0.154976	-0.270	0
1	1990	-0.270	0.816	1.572	-3.936	-0.229	0.148574	-2.576	0
2	1990	-2.576	-0.270	0.816	1.572	-3.936	0.159837	3.514	1
3	1990	3.514	-2.576	-0.270	0.816	1.572	0.161630	0.712	1
4	1990	0.712	3.514	-2.576	-0.270	0.816	0.153728	1.178	1

test.head:

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
933	2008	-4.522	-0.402	1.125	-2.440	1.588	3.372257	-0.752	0
934	2008	-0.752	-4.522	-0.402	1.125	-2.440	4.788802	-5.412	0
935	2008	-5.412	-0.752	-4.522	-0.402	1.125	5.006464	0.409	1

```

936 2008 0.409 -5.412 -0.752 -4.522 -0.402 5.100980 4.871
1
937 2008 4.871 0.409 -5.412 -0.752 -4.522 4.539542 -4.596
0

```

test obs: 156

Generalized Linear Model Regression Results

```

=====
=====
Dep. Variable:          Direction  No. Observations:
933
Model:                GLM  Df Residuals:
931
Model Family:         Binomial  Df Model:
1
Link Function:        logit  Scale:
1.0000
Method:                IRLS  Log-Likelihood:
-639.25
Date:                  Mon, 15 Feb 2021  Deviance:
1278.5
Time:                  17:17:25  Pearson chi2:
933.
No. Iterations:                4
Covariance Type:              nonrobust
=====
=====

```

	coef	std err	z	P> z	[0.025
0.975]					

Intercept	0.2266	0.066	3.422	0.001	0.097
0.356					
Lag2	0.0472	0.032	1.460	0.144	-0.016
0.110					

Xtest.head:

	Lag1	Lag2	Lag3	Lag4	Lag5	Volume
933	-4.522	-0.402	1.125	-2.440	1.588	3.372257
934	-0.752	-4.522	-0.402	1.125	-2.440	4.788802
935	-5.412	-0.752	-4.522	-0.402	1.125	5.006464
936	0.409	-5.412	-0.752	-4.522	-0.402	5.100980
937	4.871	0.409	-5.412	-0.752	-4.522	4.539542

ytest.head:

```
933    0
934    0
935    1
936    1
937    0
```

Name: Direction, dtype: int64

cm:

```
[[ 7 65]
 [ 5 79]]
```

aScore: 0.5512820512820513

```
In [94]: fraction_correct_test = fraction_correct_test
# NOTE: at a p value of 0.144, the result is not stat sig, since 0.05 <
pValue
#raise NotImplementedError
```

```
In [95]: print(f'Overall fraction of correct predictions is {fraction_correct_te
st}')
```

Overall fraction of correct predictions is 0.5512820512820513

Pass in the train and test datasets to make sure that they're working (feel free to rename the variables in the test case), and then run the test for `fraction_correct_test` !

```
In [96]: grader.grade(test_case_id = 'A5_df_test', answer = (train, test))
```

Correct! You earned 1/1 points. You are a star!

Your submission has been successfully recorded in the gradebook.

```
In [97]: grader.grade(test_case_id = 'A5_fraction_correct_test', answer = fraction_correct_test)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

Part B

Now, we want to develop an investment strategy in which we buy if the returns are greater than 0.5% and sell otherwise.

B1.

Create a response variable, y_i such that

$$y_i = \begin{cases} 1 & \text{if Today} > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

```
In [98]: encode = lambda x: 1 if x > 0.50 else 0
weekly['Response'] = weekly['Today'].map(encode)
weekly.head()
```

Out[98]:

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction	Response
0	1990	0.816	1.572	-3.936	-0.229	-3.484	0.154976	-0.270	0	0
1	1990	-0.270	0.816	1.572	-3.936	-0.229	0.148574	-2.576	0	0
2	1990	-2.576	-0.270	0.816	1.572	-3.936	0.159837	3.514	1	1
3	1990	3.514	-2.576	-0.270	0.816	1.572	0.161630	0.712	1	1
4	1990	0.712	3.514	-2.576	-0.270	0.816	0.153728	1.178	1	1


```
In [99]: grader.grade(test_case_id = 'B1_response_test', answer = weekly)
```

Correct! You earned 1/1 points. You are a star!

Your submission has been successfully recorded in the gradebook.

B2.

Fit a logistic regression model using a training data period from 1990 to 2008, with the five lag variables and volume as predictors.

```
In [100]: train = (weekly[weekly["Year"] < 2009])
#fit3 = smf.glm('y_1~Lag1+Lag2+Lag3+Lag4+Lag5+Volume', data = weekly2,
family=sm.families.Binomial()).fit()
fit3 = smf.glm('Response~Lag1+Lag2+Lag3+Lag4+Lag5+Volume', data = weekly, family=sm.families.Binomial()).fit()
```

Pass in the regression equation to `logit_equation_B` below

```
In [101]: logit_equation_B = 'Response~Lag1+Lag2+Lag3+Lag4+Lag5+Volume'
```

```
In [102]: grader.grade(test_case_id = 'B2_logit_test', answer = logit_equation_B)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

B3.

Use the `summary()` function to print the results. Do any of the predictors appear to be statistically significant? Which predictors appear to be statistically significant?

```
In [103]: #raise NotImplementedError
print(fit3.summary(), '\n')
```

```
print('Predicted probability of going Up:', fit3.predict(), '\n') #Predict
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Response   No. Observations:
1089
Model:                  GLM       Df Residuals:
1082
Model Family:          Binomial   Df Model:
6
Link Function:         logit      Scale:
1.0000
Method:                IRLS       Log-Likelihood:
-746.35
Date:                  Mon, 15 Feb 2021   Deviance:
1492.7
Time:                  17:17:42   Pearson chi2:
1.09e+03
No. Iterations:        4
Covariance Type:       nonrobust
=====
```

	coef	std err	z	P> z	[0.025
					0.975]

Intercept	-0.1616	0.086	-1.889	0.059	-0.329
0.006					
Lag1	-0.0455	0.026	-1.723	0.085	-0.097
0.006					
Lag2	0.0369	0.027	1.389	0.165	-0.015
0.089					
Lag3	-0.0128	0.026	-0.485	0.628	-0.064
0.039					
Lag4	-0.0156	0.026	-0.592	0.554	-0.067
0.036					
Lag5	-0.0142	0.026	-0.539	0.590	-0.066
0.037					

```
Volume      -0.0186      0.037      -0.504      0.614      -0.091
0.054
=====
=====
```

```
Predicted probability of going Up: [0.48991811 0.48063181 0.49083875
... 0.46030945 0.44146026 0.42579458]
```

Type the number of apparently significant variables into `num_significant_B` and the names of the variables into the list `var_significant_B` -- the test case will only give points if both variables are correct!

```
In [104]: print(fit3.pvalues)

num_significant_B = 1
var_significant_B = ['Lag1'] # This has to be a list!
#raise NotImplementedError
```

```
Intercept    0.058859
Lag1         0.084944
Lag2         0.164976
Lag3         0.627794
Lag4         0.553562
Lag5         0.590073
Volume       0.614221
dtype: float64
```

```
In [105]: grader.grade(test_case_id = 'B3_significant_test', answer = (num_signif
ificant_B, var_significant_B))
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

B4.

Compute the overall fraction of correct predictions for the held out data (that is, the data from 2009 and 2010). Assign this value to the variable `fraction_correct`. What is the overall fraction of correct predictions?

```
In [106]: from sklearn.metrics import confusion_matrix
          from sklearn.metrics import accuracy_score

          test = (weekly[weekly["Year"] >= 2009])
          Xtest = test[weekly.columns[1:-1]]
          print(Xtest.head())
          ytest = test[weekly.columns[-1]]
          print(ytest.head())
          print('len ytest: ', len(ytest), '\n')
          yprobs = fit3.predict(Xtest)
          print(yprobs)
          ypred = [0 if y < 0.5 else 1 for y in yprobs]

          x = confusion_matrix(ytest, ypred)
          print('cm: ')
          print(x, '\n')

          aScore = accuracy_score(ytest, ypred)

          print('aScore: ', aScore)
          #print('The accuracy score is', 100 * accuracy_score(data['Direction'],
          #predictions['Direction']), '%')
          fraction_correct_test = aScore

          fraction_correct = 52/104
          print("fraction_correct ", fraction_correct)
```

	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
985	6.760	-1.698	0.926	0.418	-2.251	3.793110	-4.448	0
986	-4.448	6.760	-1.698	0.926	0.418	5.043904	-4.518	0
987	-4.518	-4.448	6.760	-1.698	0.926	5.948758	-2.137	0
988	-2.137	-4.518	-4.448	6.760	-1.698	6.129763	-0.730	0
989	-0.730	-2.137	-4.518	-4.448	6.760	5.602004	5.173	1
985	0							
986	0							

```
987    0
988    0
989    1
Name: Response, dtype: int64
len ytest: 104
```

```
985    0.356914
986    0.549282
987    0.424508
988    0.408577
989    0.430381
```

```
...
1084   0.447923
1085   0.395309
1086   0.460309
1087   0.441460
1088   0.425795
Length: 104, dtype: float64
cm:
[[48  5]
 [49  2]]
```

```
aScore: 0.4807692307692308
fraction_correct 0.5
```

```
In [107]: print(f'Overall fraction of correct predictions is {fraction_correct}')
```

Overall fraction of correct predictions is 0.5

```
In [108]: grader.grade(test_case_id = 'B4_fraction_test', answer = fraction_correct)
```

Correct! You earned 2/2 points. You are a star!

Your submission has been successfully recorded in the gradebook.

Submit

You're done! Please make sure you've run all the PennGrader cells and count up your score to be sure (there are 20 points in total) and then make sure to submit this on Codio.

In []: